

# the consensus mechanism of Conflux



# a bit about me...

Péter Garamvölgyi

I did my Masters in CS at Tsinghua (2018-20)

I served as the previous president of TIBA

I joined Conflux in May 2019



scalability



decentralization

security



## P2P Network



~10Mbps  
=> 3000-6000 TPS

## Storage I/O




20K-40K SSD I/O Operations  
=> 2000-10000 TPS

## Compute



4000~8000  
Transaction Signature  
Verification  
per Second





# consensus in Bitcoin

everyone agrees which transactions happened in which order

components: **block tree**

longest chain rule

proof-of-work



# block tree in Bitcoin


transactions are packed into blocks



# block tree in Bitcoin

transactions are packed into blocks

each block references exactly one previous block






# block tree in Bitcoin

transactions are packed into blocks

each block references exactly one previous block

the chain of blocks defines the canonical transaction order




# block tree in Bitcoin

transactions are packed into blocks


each block references exactly one previous block

the chain of blocks defines the canonical transaction order

**two blocks can reference the same parent block**



# block tree in Bitcoin



# consensus in Bitcoin

everyone agrees which transactions happened in which order

components: block tree


longest chain rule

proof-of-work



# longest chain rule in Bitcoin


how do we handle concurrent blocks?



# longest chain rule in Bitcoin

how do we handle concurrent blocks?

always select the longest chain, discard the rest




# longest chain rule in Bitcoin

how do we handle concurrent blocks?

always select the longest chain, discard the rest

newly mined blocks will reference the longest chain




# longest chain rule in Bitcoin

how do we handle concurrent blocks?

always select the longest chain, discard the rest

newly mined blocks will reference the longest chain





# consensus in Bitcoin

everyone agrees which transactions happened in which order

components: block tree


longest chain rule

proof-of-work



# proof-of-work


it is trivial to make a competing long chain of blocks



# proof-of-work

it is trivial to make a competing long chain of blocks

we artificially make it hard to create blocks




# proof-of-work

it is trivial to make a competing long chain of blocks

we artificially make it hard to create blocks

an attacker with  $<50\%$  hashpower cannot keep up



# consensus in Bitcoin

everyone agrees which transactions happened in which order

components: block tree

longest chain rule

proof-of-work


problem: discarded blocks do not contribute to the security or throughput of the system.

can we do better?




# GHOST chain selection

what happens if we increase the block creation rate?




# GHOST chain selection

what happens if we increase the block creation rate?



# GHOST chain selection

what happens if we increase the block creation rate?




the longest chain has 37.5%  
of the network's hashpower

all transactions from 5  
blocks are discarded




# GHOST chain selection

instead of the longest chain rule, we will use the GHOST rule




# GHOST chain selection

instead of the **longest chain rule**, we will use the GHOST rule




# GHOST chain selection

instead of the **longest chain rule**, we will use the **GHOST** rule



# GHOST chain selection

instead of the **longest chain rule**, we will use the **GHOST** rule



# GHOST chain selection

instead of the longest chain rule, we will use the GHOST rule

with GHOST, we always choose the subtree of the majority

this makes GHOST more resilient, even under high block rates




# GHOST chain selection

instead of the longest chain rule, we will use the GHOST rule  
with GHOST, we always choose the subtree of the majority  
this makes GHOST more resilient, event under high block rates  
can GHOST still be attacked?




# liveness attack on GHOST

an attacker pre-mine blocks and maintain long forks



# liveness attack on GHOST


an attacker pre-mine blocks and maintain long forks






# liveness attack on GHOST

an attacker pre-mine blocks and maintain long forks




# liveness attack on GHOST

an attacker pre-mine blocks and maintain long forks



# liveness attack on GHOST


an attacker pre-mine blocks and maintain long forks



# GHOST chain selection

an attacker pre-mine blocks and maintain long forks

when we detect a liveness attack, we switch to adaptive mode




# GHOST chain selection

an attacker pre-mine blocks and maintain long forks

when we detect a liveness attack, we switch to adaptive mode

occasional rare blocks will break the balance



# GHASt chain selection

an attacker pre-mine blocks and maintain long forks

when we detect a liveness attack, we switch to adaptive mode

occasional rare blocks will break the balance

**GHASt: Greedy Heaviest Adaptive Sub-Tree**



so where are we now?

with GHOST, the chain is stable even under high block rates

with GHAST, the chain is resilient against liveness attacks



so where are we now?

with GHOST, the chain is stable even under high block rates

with GHOST, the chain is resilient against liveness attacks

but discarded blocks still do not contribute to throughput


is there a way we could use them?






# tree-graph

let's look at a simple fork




# tree-graph

we know the A comes before C, but B and C are concurrent



# tree-graph


let's introduce reference edges that capture the order: B-C




# tree-graph

let's introduce reference edges that capture the order: B-C

this additional information will help us derive an order




tree-graph




# tree-graph

parent edges form a tree




# tree-graph

all edges form a DAG



# tree-graph


we perform GHA<sup>ST</sup> on the parent-tree to get the pivot chain





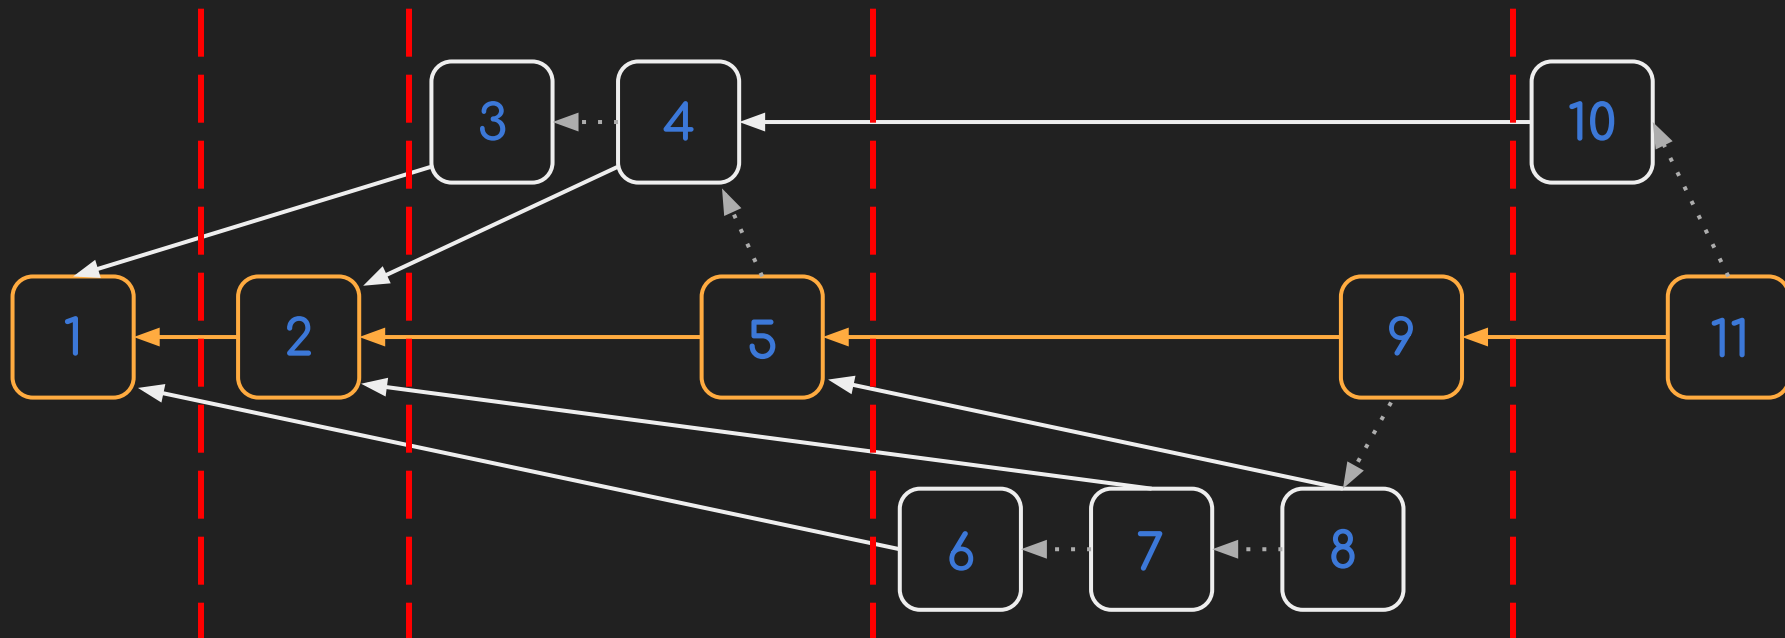
# block order in the tree-graph

then, we divide the graph into epochs




# block order in the tree-graph

we can derive a linear order of blocks from the epochs




# transaction order in the tree-graph

from a tree graph, we can also derive the transaction order



# transaction order in the tree-graph



tx-A tx-B tx-C tx-D tx-E tx-C tx-F tx-G tx-F tx-E tx-H



# things we learned today

Bitcoin uses Proof-of-Work and the longest chain rule to achieve consensus on a linear sequence of transactions



# things we learned today

Bitcoin uses Proof-of-Work and the longest chain rule to achieve consensus on a linear sequence of transactions

with GHOST, the main chain remains stable, even with many forks, and it is resilient against liveness attacks



# things we learned today

Bitcoin uses Proof-of-Work and the longest chain rule to achieve consensus on a linear sequence of transactions

with GHOST, the main chain remains stable, even with many forks, and it is resilient against liveness attacks

with the Tree-Graph ledger structure, we can keep transactions from all blocks, thus increasing the overall throughput



thank you!





# resources

Sompolinsky, Yonatan, and Aviv Zohar. "Secure high-rate transaction processing in bitcoin." 2015.

<https://eprint.iacr.org/2013/881.pdf>

Li, Chenxin, Peilun Li, Dong Zhou, Zhe Yang, Ming Wu, Guang Yang, Wei Xu, Fan Long, and Andrew Chi-Chih Yao. "A decentralized blockchain with high throughput and fast confirmation." 2020.

<https://www.usenix.org/conference/atc20/presentation/li-chenxing>

Conflux technical presentation

[https://confluxnetwork.org/files/Conflux\\_Technical\\_Presentation\\_20200309.pdf](https://confluxnetwork.org/files/Conflux_Technical_Presentation_20200309.pdf)

